

Getting started with LinuxBoot Firmware on AArch64 Server

Naohiro Tamura (naohiro.tamura@fujitsu.com)

Professional Engineer
Fujitsu Limited

Slides: <https://github.com/NaohiroTamura/LCA2021>

Outline

- My Motivation, Your Merit, and Our Goal
- What is LinuxBoot and its 2 Pitfalls?

- Solution
- How to create, boot and debug Flashrom
 - Tip 1: Create AArch64 OVMF 32MB Firmware File System
 - Tip 2: Configure LinuxBoot Kernel and Initramfs
 - Tip 3: Inject LinuxBoot into QEMU 64MB Flashrom
 - Tip 4: Boot Final OS from Local Disk
 - Tip 5: Debug LinuxBoot AArch64 Kernel using QEMU and GDB on x86_64

- What To Do Next?
- Summary

Key Words

2 Pitfalls



1 Solution



5 Tips

My Motivation, Your Merit and Our Goal



■ My Motivation

■ Don't Repeat My Struggle by sharing **2 Pitfalls, 1 Solution and 5 Tips.**

- Last year I investigated LinuxBoot for AArch64 Server Project.
- Because **LinuxBoot is Mega Datacenter Customer's Requirement for Security.**



■ Your Merit

■ Be able to explain LinuxBoot AArch64 to your Boss and Colleagues with 100% Confidence.

- Because of getting LinuxBoot AArch64 Box Today without purchasing any additional HW at all.

■ Our Goal

■ Boot Final OS, CentOS 8.2 AArch64, from LinuxBoot Flashrom using QEMU by Ourselves

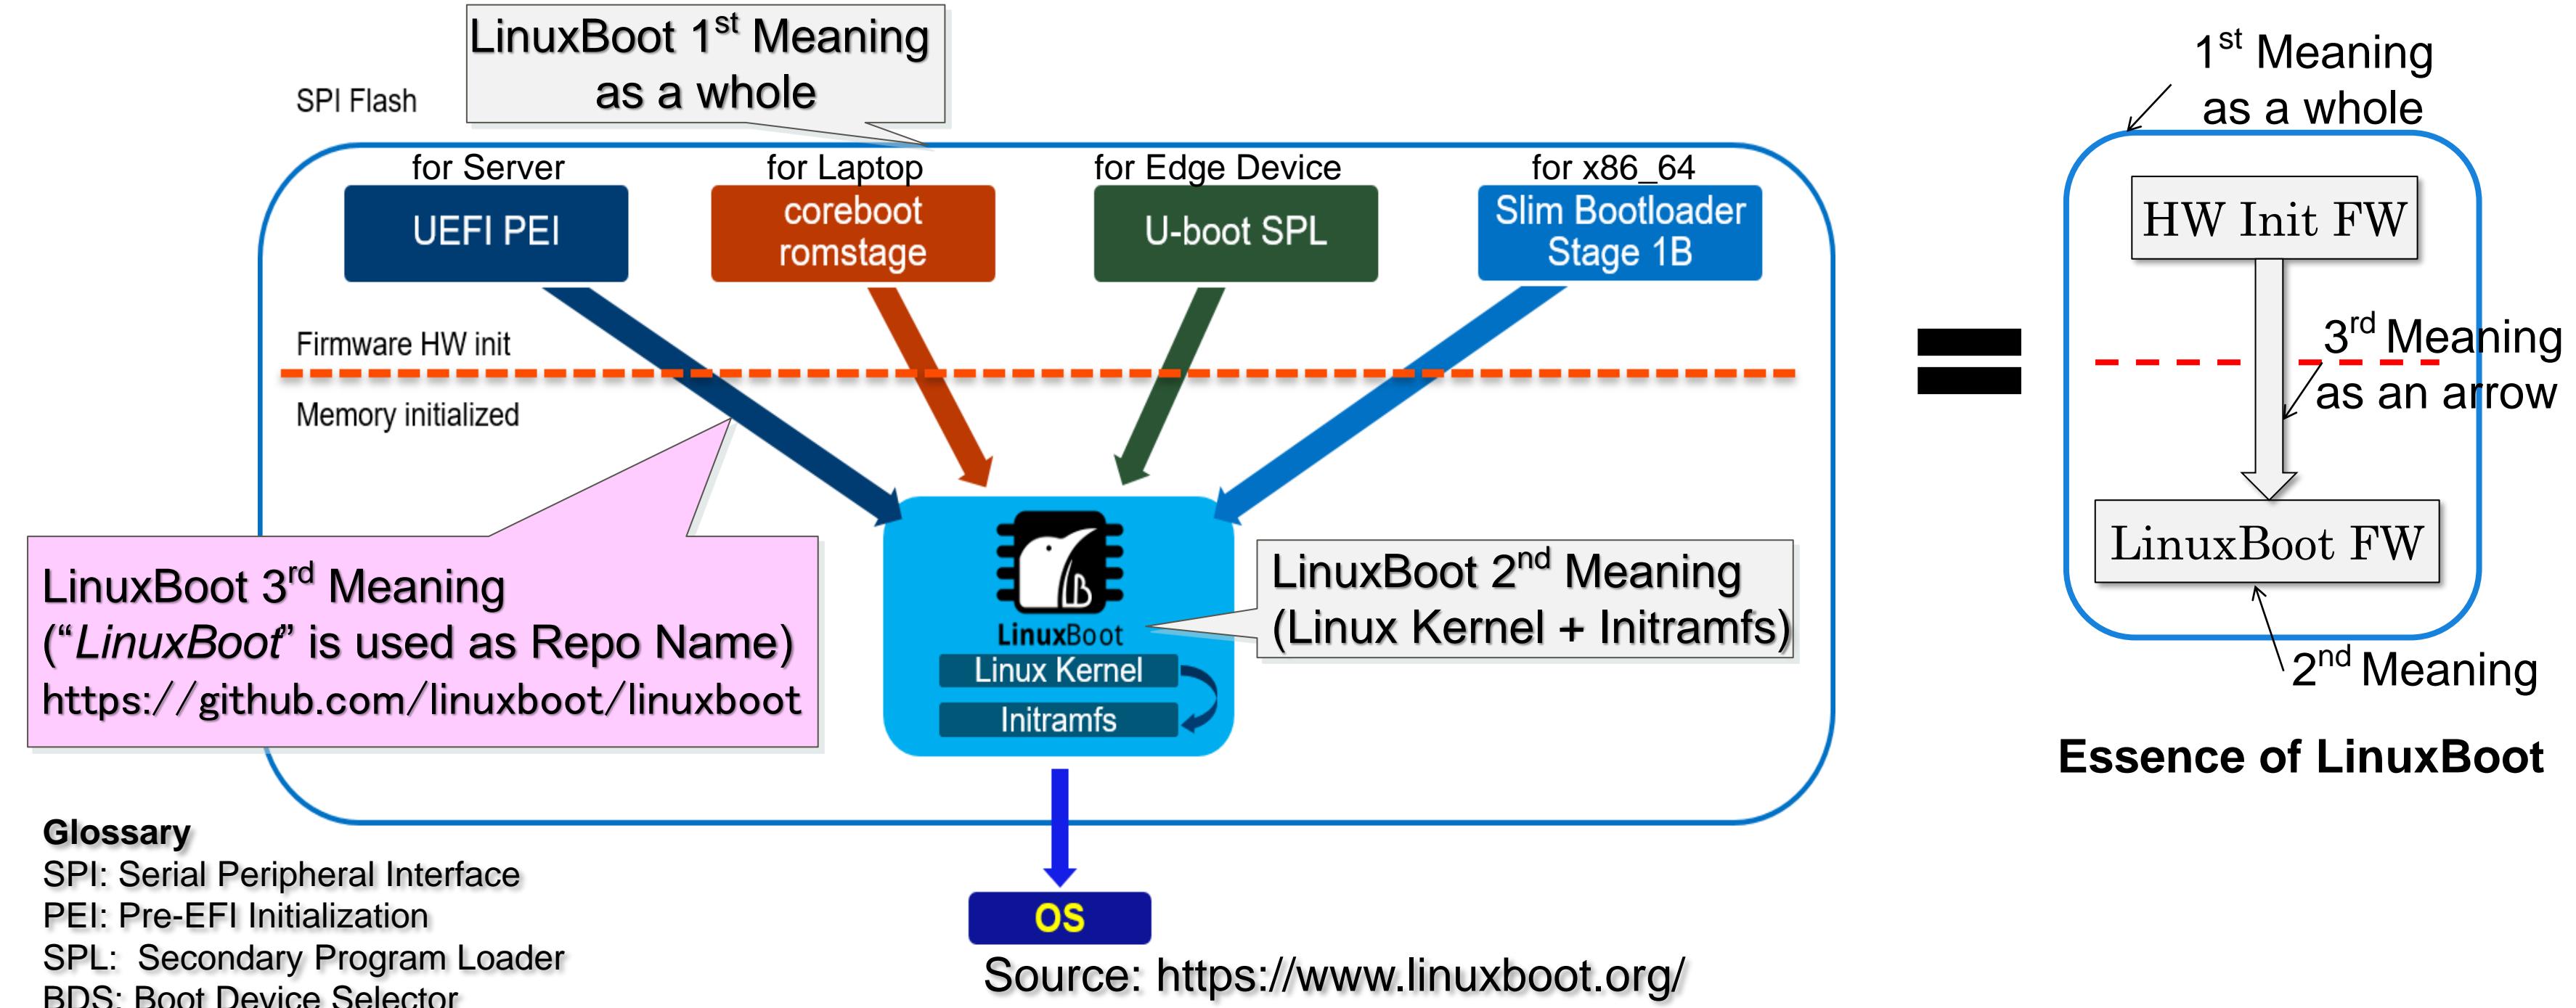
- Consider CentOS 8.2 as RHEL 8.2 which enterprise customers mostly use
- All steps to reproduce are available at <https://github.com/NaohiroTamura/LCA2021>

* Source: <https://www.bloomberg.com/news/articles/2020-12-18/microsoft-is-designing-its-own-chips-for-servers-surface-pcs>

What is LinuxBoot?

■ “LinuxBoot” has Three Meanings depending on Contexts.

■ We focus on LinuxBoot 3rd Meaning (UEFI PEI to LinuxBoot 2nd) because it's for Server



The 1st Pitfall: No BDS Kernel Param and Patch for AArch64 FUJITSU

■ No such Kernel Param **CONFIG_EFI_BDS**

- The [GitHub](#) provides no further instructions

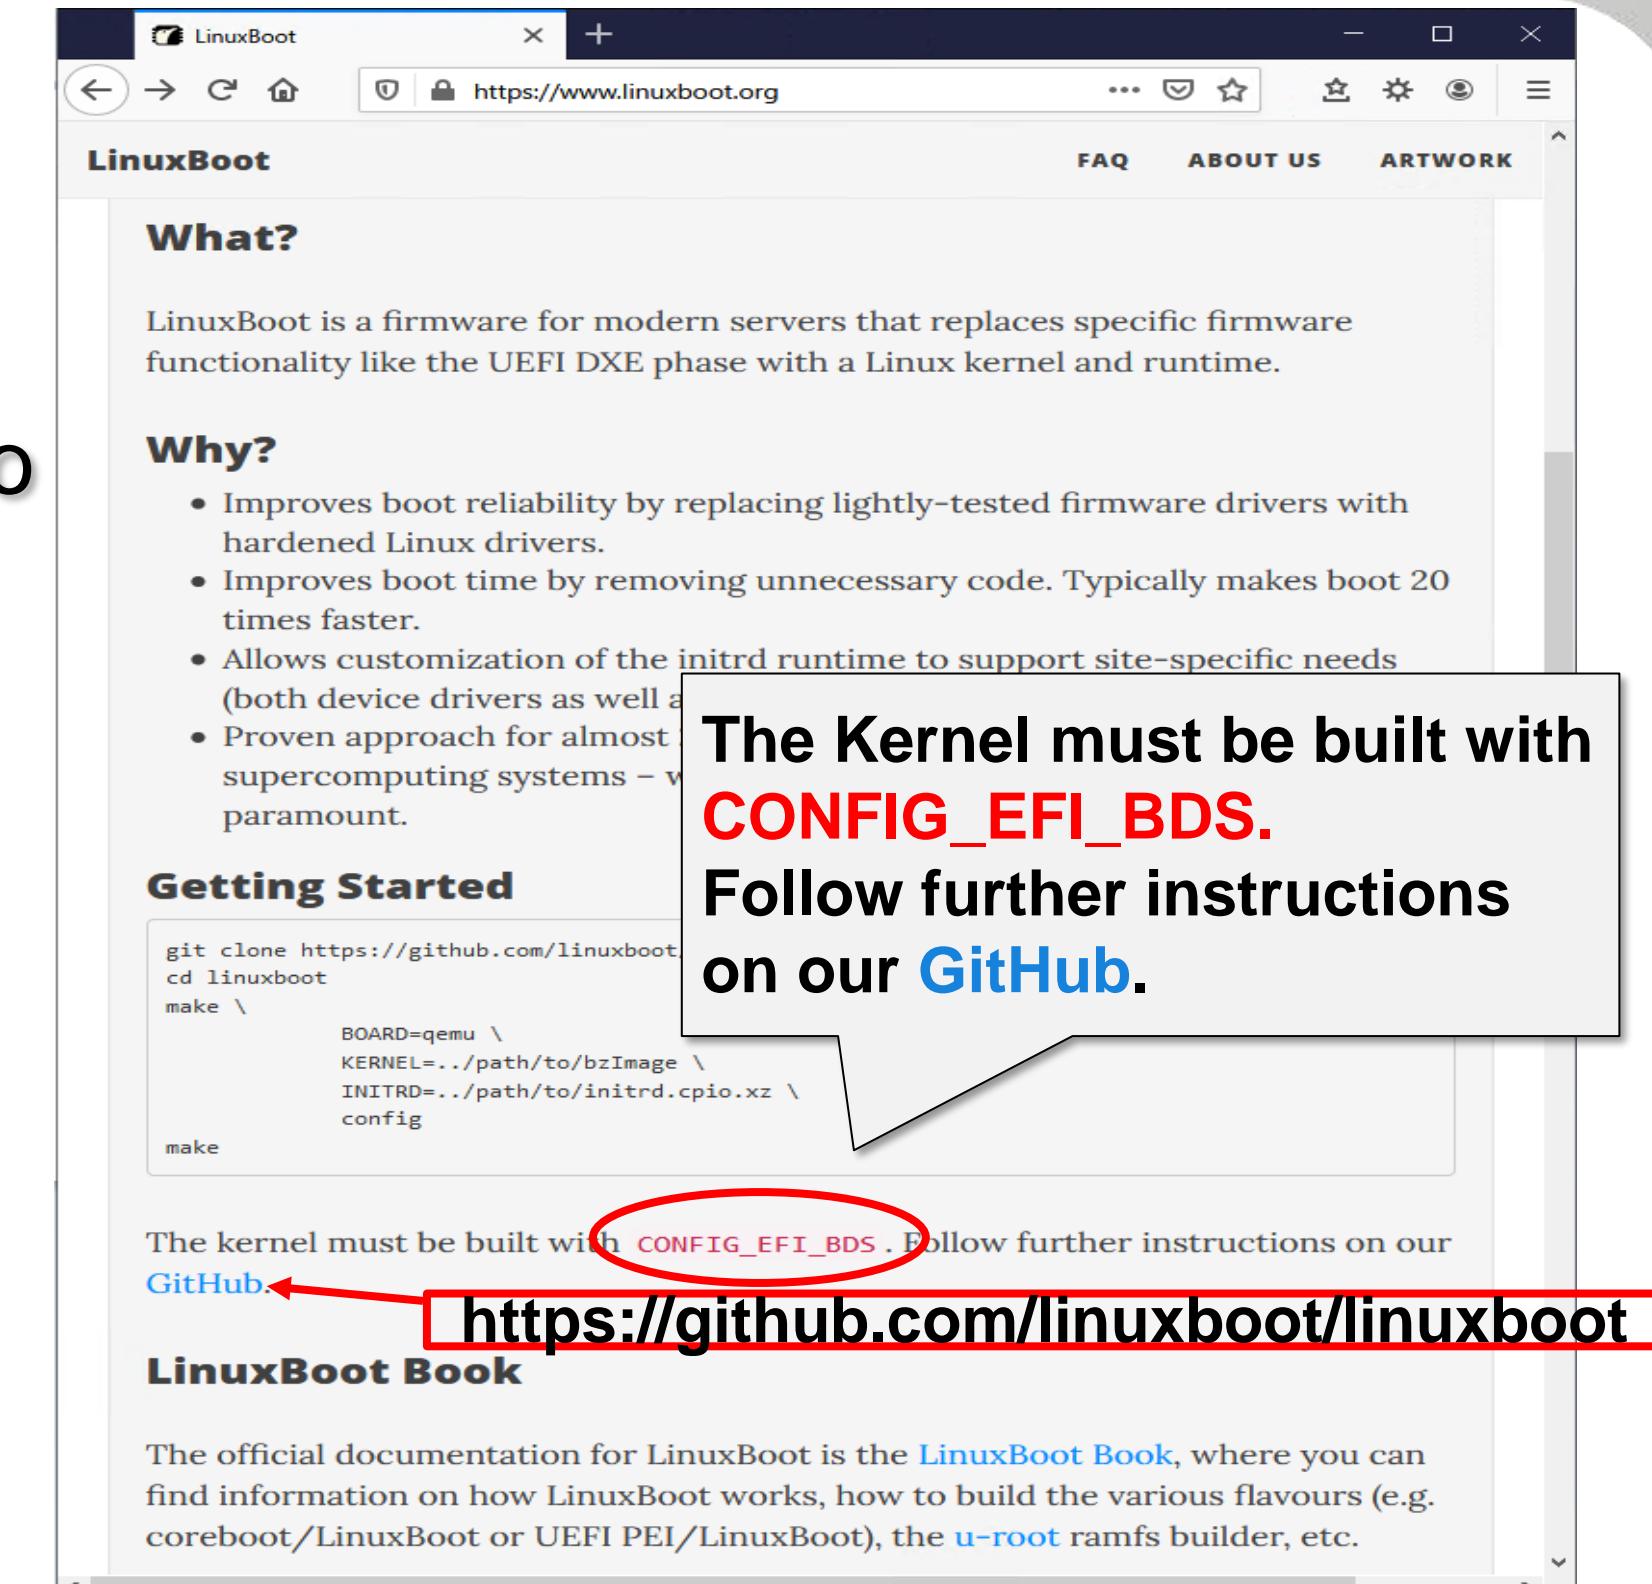
■ But found BDS Kernel Patch in HEADS repo

- This Kernel Patch is to x86_64 arch dependent code, so **it's NOT applicable to AArch64**

- https://github.com/osresearch/heads/blob/master/patches/linux-4.14.62/0000-efi_bds.patch

■ Basically What is BDS ?

- BDS (Boot Device Selector)



Source: <https://www.linuxboot.org/>

The 2nd Pitfall: No LinuxBoot BDS Source Code for AArch64 FUJITSU

- BDS is a phase of UEFI Boot. LinuxBoot BDS selects Flashrom Device and boot
 - <https://github.com/linuxboot/linuxboot/blob/master/dxe/linuxboot.c>

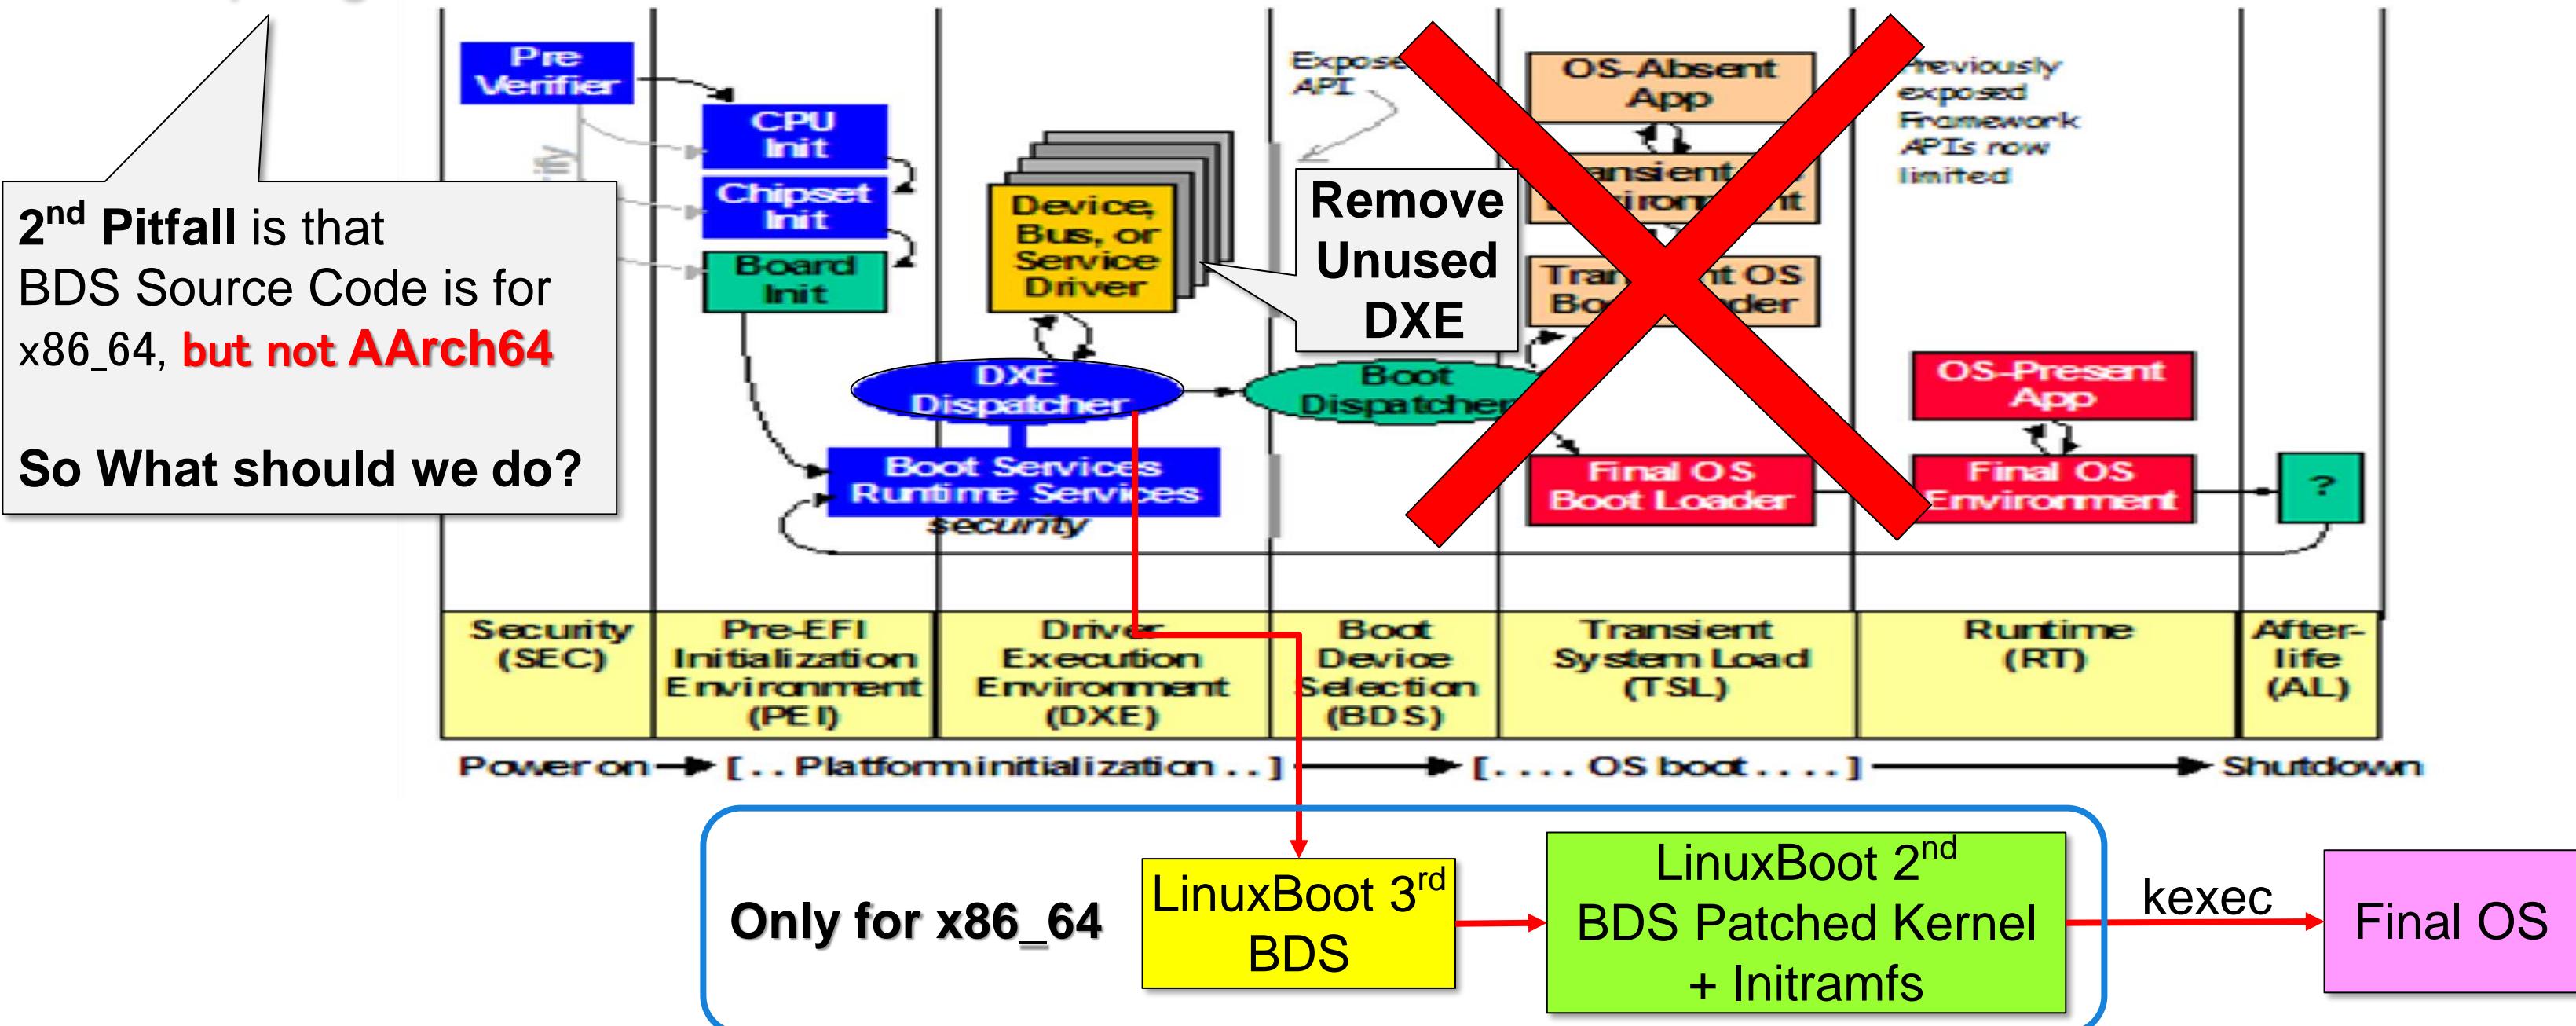
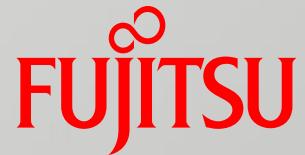


Figure Source: https://edk2-docs.gitbook.io/edk-ii-build-specification/2_design_discussion/23_boot_sequence

Solution: Replace UEFI Shell with LinuxBoot in Flashrom



■ Fiano replaces UEFI Shell with LinuxBoot then BootManager calls LinuxBoot 2nd

■ <https://github.com/linuxboot/fiano>

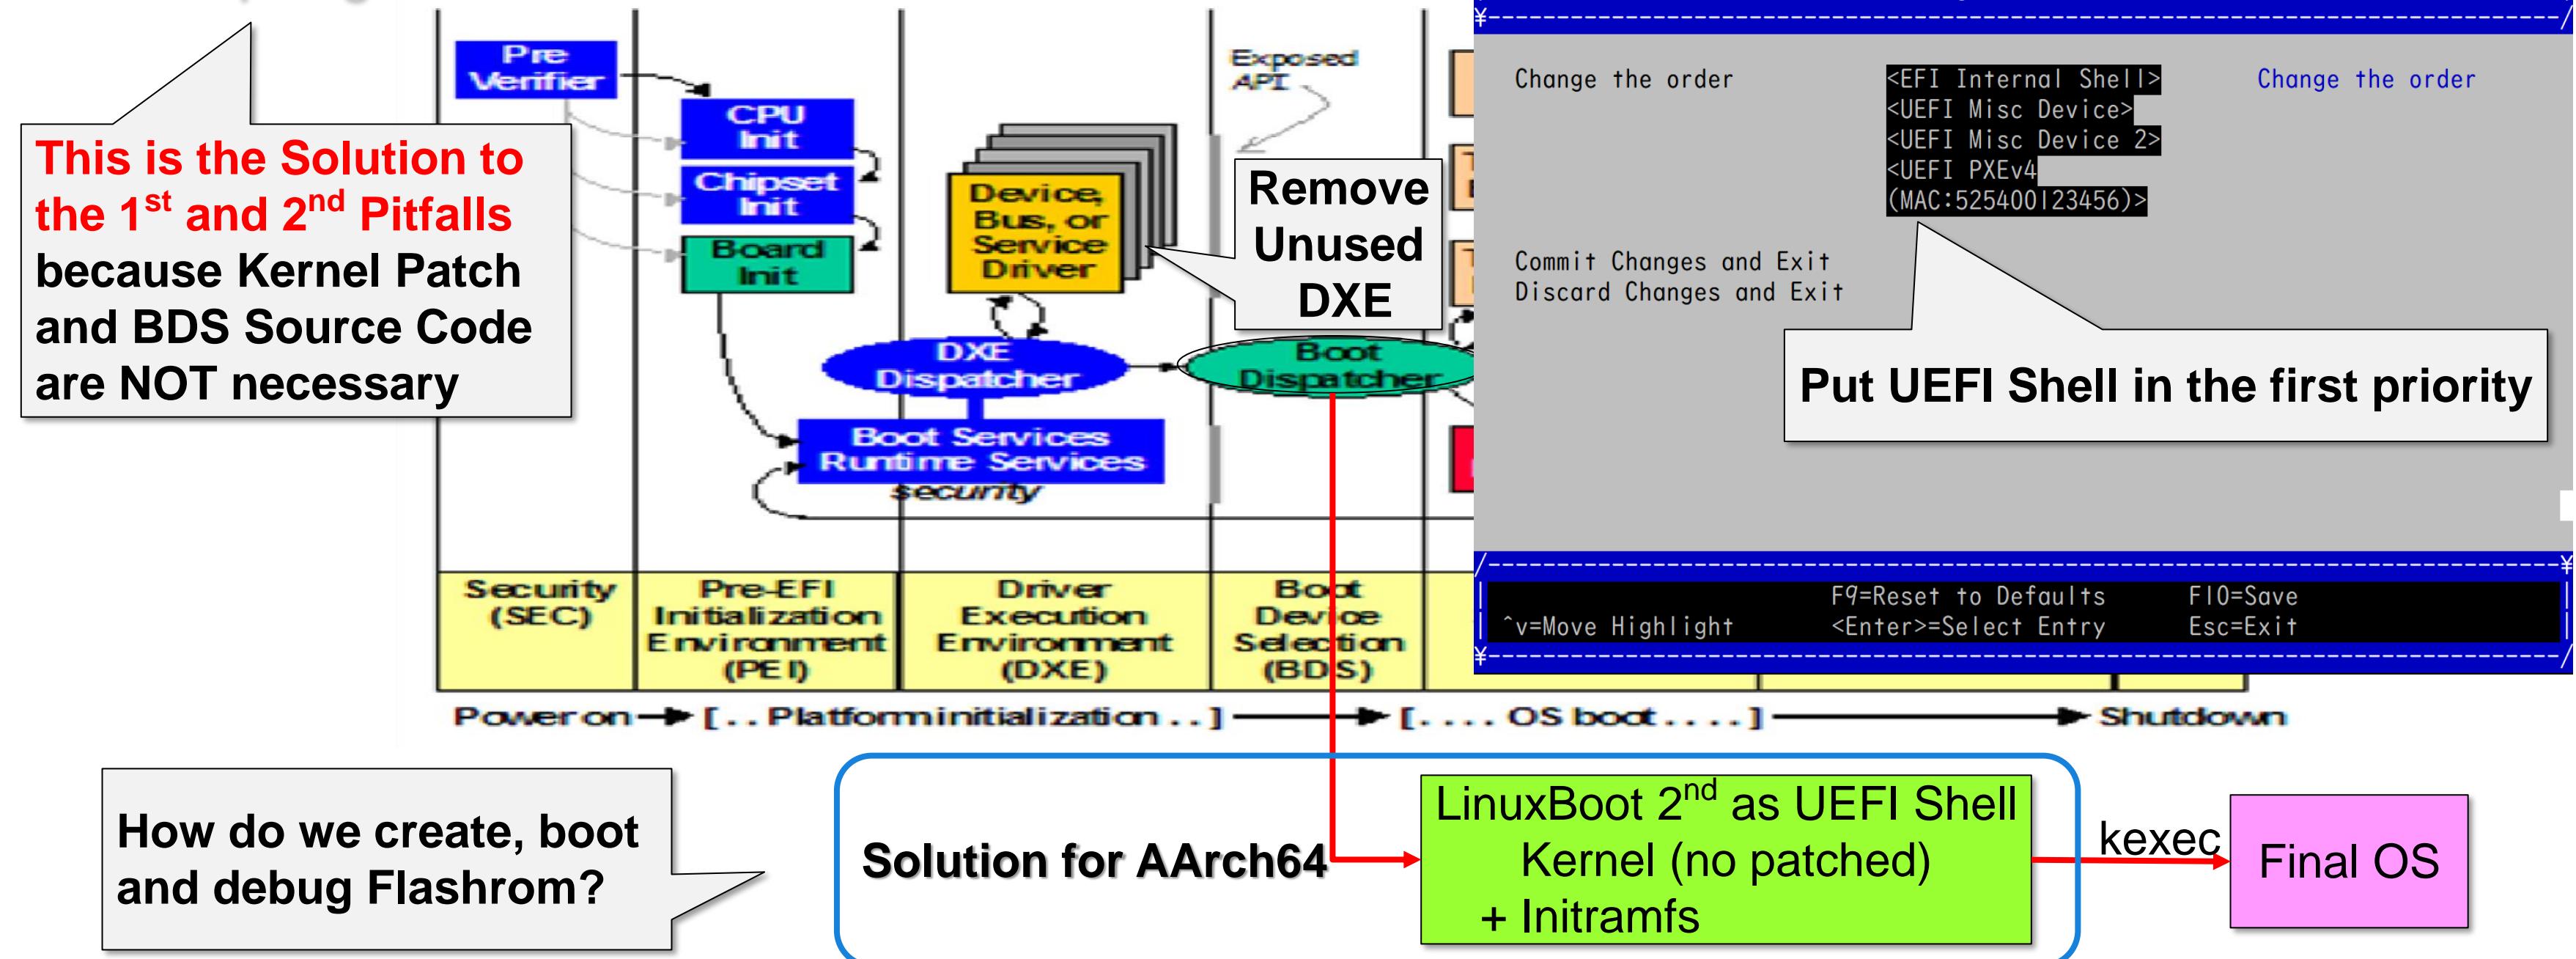
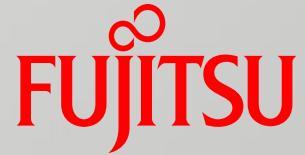


Figure Source: https://edk2-docs.gitbook.io/edk-ii-build-specification/2_design_discussion/23_boot_sequence

5 Tips for LinuxBoot AArch64



■ How to create, boot and debug Flashrom

- Tip 1: Create AArch64 OVMF 32MB Firmware File System
- Tip 2: Configure LinuxBoot Kernel and Initramfs
- Tip 3: Inject LinuxBoot into QEMU 64MB Flashrom
- Tip 4: Boot Final OS from Local Disk
- Tip 5: Debug LinuxBoot AArch64 Kernel using QEMU and GDB on x86_64

■ Flashrom Size Requirement and Challenge

■ Flashrom Size is 32MB

- Low End Physical AArch64 Server has only 32MB Flashrom.
- Trusted Firmware (8MB)+ UEFI(8MB) + LinuxBoot (Kernel + Initramfs) < 32MB

LinuxBoot 2nd (Kernel + Initramfs) has to be less than 16MB without Kernel Compression

■ AArch64 Kernel has to be stored in uncompressed (3 times larger than compressed)

- Because AArch64 doesn't support Self Decompression PE/COFF Kernel Image, but X86, x86_64 and AArch32 do
- FYI, CentOS 8.2 generic kernel size gzip 8MB and gunzip 25MB

Tip 1: Create AArch64 OVMF 32MB Firmware File System

■ OVMF File Size is only 2MB, so no room to replace UEFI Shell with LinuxBoot 2nd

■ OVMF (Open Virtual Machine Firmware) is UEFI implementation for QEMU and KVM

- AArch64 <https://github.com/tianocore/edk2/tree/master/ArmVirtPkg>

■ How to extend Firmware File System to 32MB?

■ Increase Flash Device # of Blocks in OVMF Source Code

- FD Block Size = 4096 Byte
- FD Size 2MB = 512 Blocks ↗
- FD Size 32MB = 8,192 Blocks ↗

Apply Patch and Rebuild OVMF

- <https://github.com/NaohiroTamura/edk2/compare/edk2-stable202008...aarch64-flashrom.patch>
ArmVirtPkg/ArmVirt.dsc.inc | 7 ++++++
ArmVirtPkg/ArmVirtQemu.fdf | 4 ++++
2 files changed, 10 insertions(+), 1 deletion(-)

Tip 2: Configure LinuxBoot Kernel and Initramfs



■ LinuxBoot 2nd (Kernel + Initramfs) has to be One File, and Size < 16MB

■ How to Minimize Kernel with embedded Initramfs?

■ Repeat Kernel Config Trial and Error using GDB

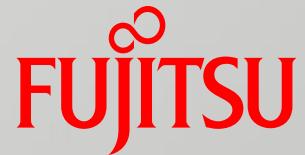
- CONFIG_EFI_STUB=y
- CONFIG_INITRAMFS_SOURCE="../initramfs.linux_arm64.cpio.xz" ←
- CONFIG_INITRAMFS_COMPRESSION_XZ=y
- # CONFIG_MODULES is not set
- Enable ACPI Support
- Minimized Kernel Defconfig is available (https://github.com/NaohiroTamura/LCA2021/blob/master/linuxboot-5.9.0-aarch64_defconfig)

→ ■ Chose u-root as Initramfs

- u-root is implemented in Golang for Security (<https://github.com/u-root/u-root>)
- Build with minimum commands
\$ GOARCH=arm64 u-root -build=bb -o=initramfs.linux_arm64.cpio -uinitcmd=boot core github.com/u-root/u-root/cmds/boot/boot
- XZ compress to 3.5MB
\$ xz --check=crc32 -9 --lzma2=dict=1MiB --stdout initramfs.linux_arm64.cpio | dd conv=sync bs=512 of=initramfs.linux_arm64.cpio.xz

LinuxBoot 2nd became 15MB
(Kernel 5.9 with embedded u-root)

Tip 3: Inject LinuxBoot into QEMU 64MB Flashrom



■ QEMU ‘virt’ machine requires 64MB Flashrom, but not 32MB

■ How to replace UEFI Shell with LinuxBoot 2nd?

■ First use ‘dd’ to extend 32MB Flashrom to 64MB by just filling out Zero

- OVMF RPM Spec file https://src.fedoraproject.org/rpms/edk2/blob/master/f/edk2.spec#_384-386
\$ dd of="arm/QEMU_EFI-pflash.raw" if="/dev/zero" bs=1M count=64
\$ dd of="arm/QEMU_EFI-pflash.raw" if="arm/QEMU_EFI.fd" conv=notrunc

■ Then use ‘replace_pe32’ subcommand of Fiano ‘utk’

- <https://github.com/linuxboot/fiano>

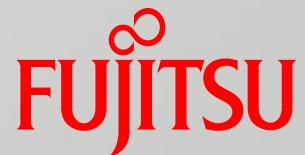
LinuxBoot 2nd PE/COFF Image 15MB
(Kernel 5.9 with embedded u-root)

Input: OVMF Flashrom 64MB

```
$ utk QEMU_EFI-pflash.raw replace_pe32 Shell build-5.9.15/arch/arm64/boot/Image ¥  
> save QEMU_EFI-pflash-linux.raw
```

Output: Linux Boot Flashrom 64MB

Tip 4: Boot Final OS from Local Disk



■ CentOS 8 follows Boot Loader Spec that u-root (Initramfs) hasn't implemented yet

■ Boot Configuration Format, Grub2 'menuentry', is changed

- https://systemd.io/BOOT_LOADER_SPECIFICATION/

■ How to boot Final OS, CentOS 8.2 from Local Disk using QEMU?

■ Apply Quick Hack Patch to u-root and rebuild LinuxBoot Flashrom

- <https://github.com/NaohiroTamura/u-root/compare/04f343dd1922457c530a90b566789fe1707d591d...centos8-bls-support.patch>

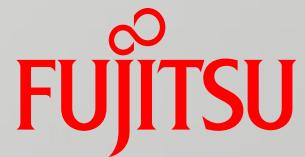
```
$ /opt/qemu-5.1.0/bin/qemu-system-aarch64 -m 8192  
-drive if=pflash,format=raw,readonly,file=QEMU_EFI-pflash-linux.raw  
-drive if=pflash,format=raw,file=vars-template-pflash.raw  
-device virtio-rng-pci -nographic -serial mon:stdio  
-machine virt,accel=tcg -cpu cortex-a72 -smp 4  
-hda centos8-aarch64-lvm.qcow2
```

Final OS CentOS 8.2

LinuxBoot Flashrom
Linux Kernel 5.9 (embedded u-root Initramfs)

'virt' machine supports ACPI

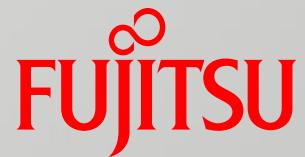
Tip 4: Boot Final OS from Local Disk (Console Log)



```
Memory Type Information settings change.  
[Bds]Booting EFI Internal Shell  
[Bds] Expand Fv(64074AFE-340A-4BE6-94BA-91B5B4D0F71E)/FvFile(7C04A583-9E3E-4F1C-AD65-E05268D0B4D1) -> Fv(64074AFE-340A-4BE6-94BA-91B5B4D0F71E)/FvFile(7C04A583-9E3E-4F1C-AD65-E05268D0B4D1)  
BdsDxe: loading Boot0002 "EFI Internal Shell" from Fv(64074AFE-340A-4BE6-94BA-91B5B4D0F71E)/FvFile(7C04A583-9E3E-4F1C-AD65-E05268D0B4D1)  
InstallProtocolInterface: 5B1B31A1-9562-11D2-8E3F-00A0C969723B 23AAF5440  
Loading driver at 0x00235AE0000 EntryPoint=0x002364AF9BC  
Loading driver at 0x00235AE0000 EntryPoint=0x002364AF9BC  
InstallProtocolInterface: BC62157E-3E33-4FEC-9920-2D3B36D750DF 23AAF2718  
ProtectUefiImageCommon - 0x3AAF5440  
- 0x0000000235AE0000 - 0x0000000000F70000  
SetUefiImageMemoryAttributes - 0x0000000235AE0000 - 0x0000000000010000 (0x0000000000004008)  
SetUefiImageMemoryAttributes - 0x0000000235AF0000 - 0x000000000009F0000 (0x00000000000020008)  
SetUefiImageMemoryAttributes - 0x00000002364E0000 - 0x00000000000570000 (0x0000000000004008)  
BdsDxe: starting Boot0002 "EFI Internal Shell" from Fv(64074AFE-340A-4BE6-94BA-91B5B4D0F71E)/FvFile(7C04A583-9E3E-4F1C-AD65-E05268D0B4D1)  
EFI stub: Booting Linux Kernel...  
EFI stub: Generating empty DTB  
EFI stub: Exiting boot services and installing virtual address map...  
SetUefiImageMemoryAttributes - 0x000000023BE80000 - 0x0000000000040000 (0x0000000000000008)
```

```
Booting Linux on physical CPU 0x00000000000 [0x410fd083]  
Linux version 5.9.15 (ubuntu@bionic) (aarch64-linux-gnu-gcc (Ubuntu/Linaro 7.5.0-3ubuntu1~18.04) 7.5.0, GNU ld  
(GNU Binutils for Ubuntu) 2.30) #1 SMP Tue Dec 22 11:18:00 UTC 2020  
efi: EFI v2.70 by EDK II  
efi: SMBIOS 3.0=0x23bef0000 MEMATTR=0x239488698 ACPI 2.0=0x238830000 RNG=0x23bffcd98 MEMRESERVE=0x238b63f18  
efi: seeding entropy pool
```

Tip 5: Debug LinuxBoot AArch64 using QEMU and GDB



■ Terminal 1

```
$ /opt/qemu-5.1.0/bin/qemu-system-aarch64 -s -S -m 8192  
-drive if=pflash,format=raw,readonly,file=QEMU_EFI-pflash-linux.raw  
-drive if=pflash,format=raw,file=vars-template-pflash.raw  
-device virtio-rng-pci -nographic -serial mon:stdio  
-machine virt,accel=tcg -cpu cortex-a72  
-hda centos8-aarch64-lvm.qcow2
```

When LinuxBoot Kernel doesn't start,
GDB debug helps us find missing driver.

■ Terminal 2

```
$ /opt/gdb-9.2/bin/aarch64-gnu-linux-gdb build-5.9.15/vmlinux
```

...

Reading symbols from build-5.9.15/vmlinux...

```
(gdb) target remote :1234
```

Default Port 1234

Remote debugging using :1234

0x0000000000000000 in ?? ()

```
(gdb) b start_kernel
```

Breakpoint 1 at 0xfffffe0010990da4: file /home/ubuntu/LCA2021/linux-5.9.15/init/main.c, line 847.

```
(gdb) c
```

Continuing.

Breakpoint 1, start_kernel () at /home/ubuntu/LCA2021/linux-5.9.15/init/main.c:847

```
847 {
```

```
(gdb)
```

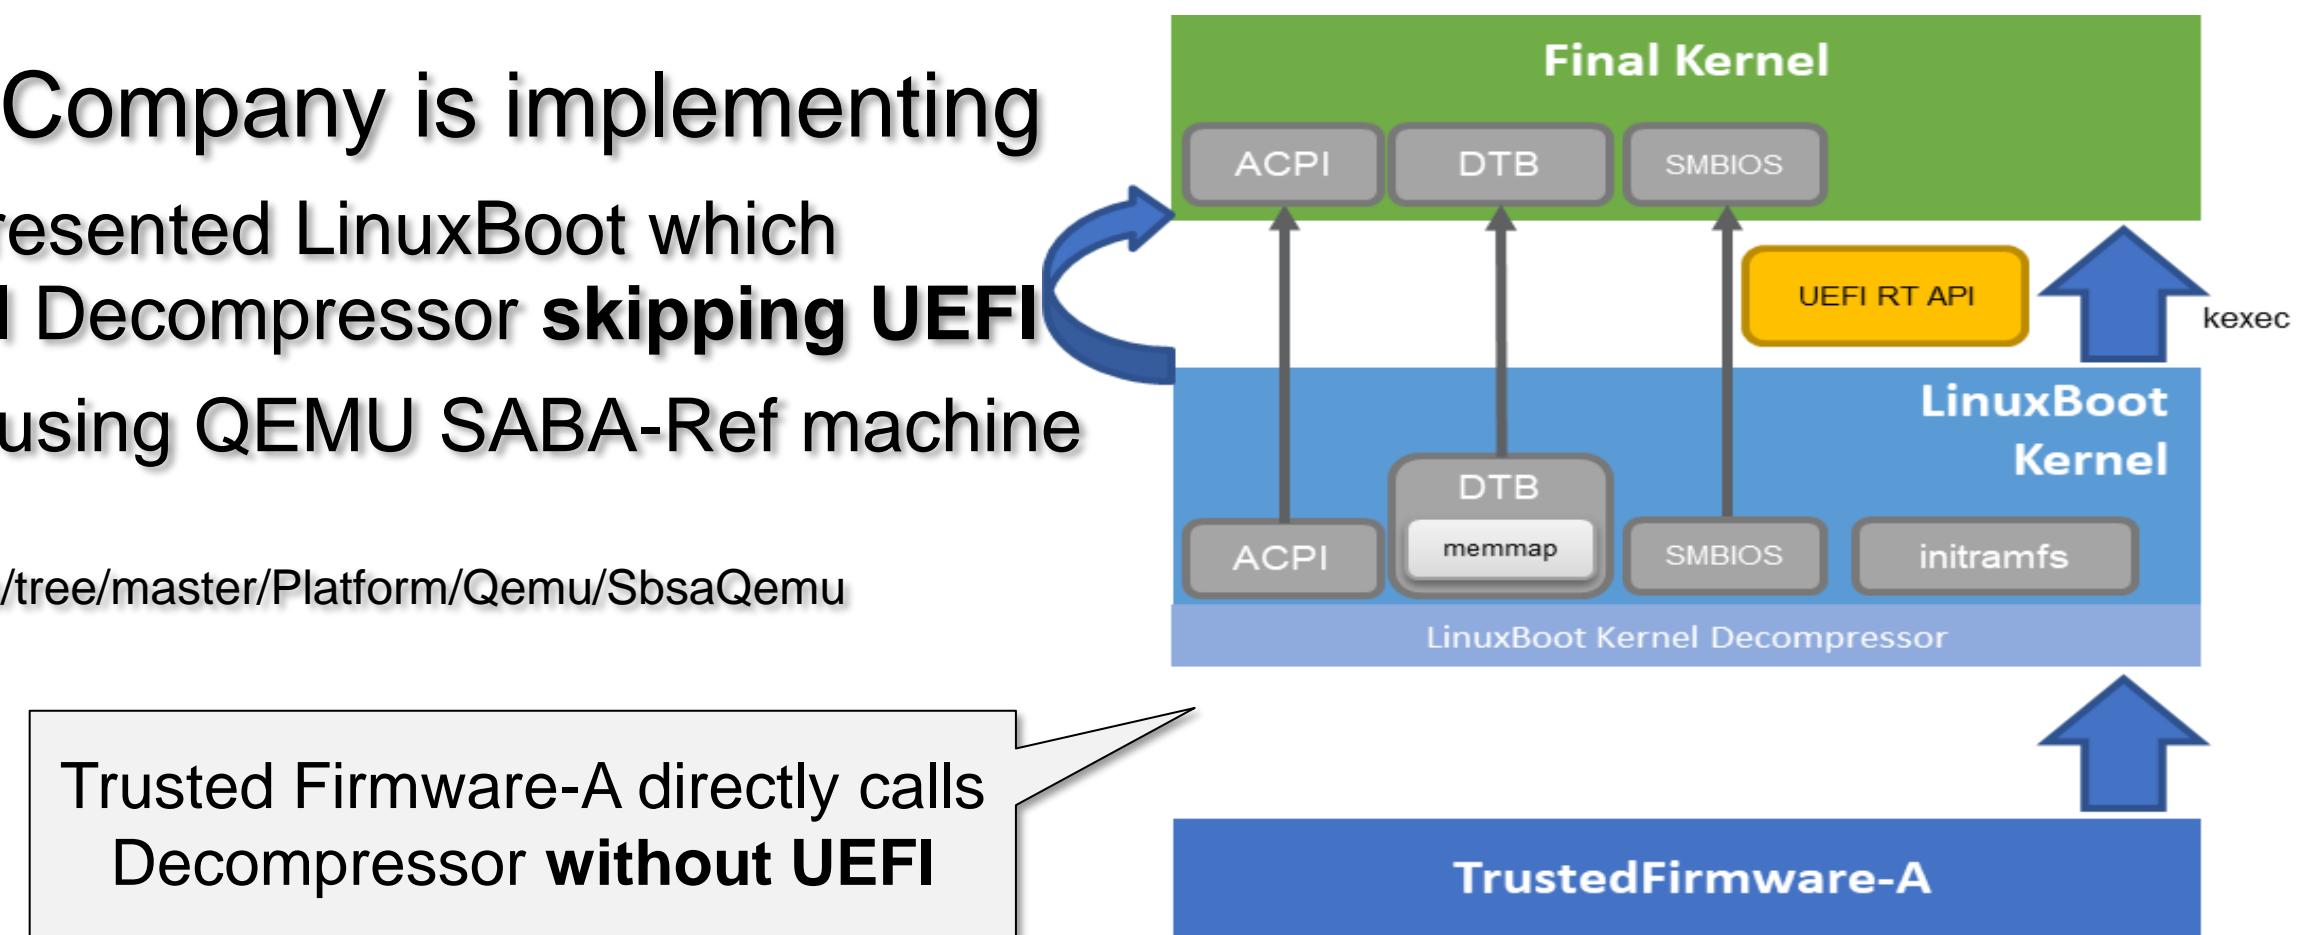
What TO DO Next?

■ Develop Kernel Decompressor UEFI Application for Fiano

- Fiano replaces UEFI Shell with the Decompressor, then the Decompressor calls LinuxBoot 2nd
- It's really peculiar why NOT only AArch64 kernel self-decompressor implemented
 - Because each loader such as Grub2, u-root and etc has to implement decompressor repeatedly
 - Found discussion once on the mailing list in Jan 2014 (<http://lists.infradead.org/pipermail/linux-arm-kernel/2014-January/224746.html>), but no more

■ Watch LinuxBoot 3rd ARM Company is implementing

- At OCP Summit 2020 ARM presented LinuxBoot which Trusted Firmware calls Kernel Decompressor **skipping UEFI**
- When it's ready, we can try it using QEMU SABA-Ref machine
 - SABA (Server Base System Architecture)
 - <https://github.com/tianocore/edk2-platforms/tree/master/Platform/Qemu/SbsaQemu>



Source: <https://2020ocpvirtualsummit.sched.com/event/bXVn/open-system-firmware-on-arm>

- Explained **2** Pitfalls, **1** Solution and **5** Tips.
- You can boot Final OS, CentOS 8.2, from LinuxBoot Flashrom using QEMU
 - All steps to reproduce are available at <https://github.com/NaohiroTamura/LCA2021>
- Try it by yourself and explain LinuxBoot AArch64 to your Boss and Colleagues
 - Please send me an email or submit an Issue to the GitHub if you had any problem.

References

■ Arm SystemReady and the UEFI firmware ecosystem

- <https://cfp.osfc.io/osfc2020/talk/KB3H9V/>

■ Open System Firmware on Arm *

- <https://2020ocpvirtualsummit.sched.com/event/bXVn/open-system-firmware-on-arm>

■ Go Forth and Modify: Fiano *

- <https://2020ocpvirtualsummit.sched.com/event/bXWK/go-forth-and-modify-fiano>

■ Firmware security, why it matters and how you can have it

- <https://2019.linux.conf.au/schedule/presentation/110/>

■ EDKII OVMF AArch64

- <https://github.com/tianocore/edk2/tree/master/ArmVirtPkg>

■ u-root

- <https://github.com/u-root/u-root>

■ LinuxBoot

- <https://github.com/linuxboot/linuxboot>

■ fiano

- <https://github.com/linuxboot/fiano>

*) Downloading slide needs to enter the OCP Virtual Summit from
<https://www.opencompute.org/summit/virtual-summit>

FUJITSU

shaping tomorrow with you